

# DEBUGGING EMBEDDED LINUX KERNEL & DRIVER

Kishore Kumar

KERNEL MASTERS [www.kernelmasters.org](http://www.kernelmasters.org)

## Contents

<b>1. KDB &amp; KGDB .....</b>	3
<b>Setup Kernel Debug Environment for KDB &amp; KGDB.....</b>	3
Step 1: Enable KGDB/KDB options.....	3
Step 2: Kernel configuration & build.....	4
Step 3: Kernel Installation.....	4
Step 4: Setup kernel bootargs in u-boot.....	4
<b>KDB Testing:.....</b>	4
Choose KDB/KGDB Enabled Kernel at bootloader:.....	4
Entering a Kernel Debugger:.....	4
<b>kdb commands .....</b>	5
<b>KGDB Debugging (source level debugging): .....</b>	6
Step 1: Connect target board remote using ssh. ....	6
Step 2: Switch KDB to KGDB.....	6
Step 3: Remote Debugging with GDB .....	6
Step 4: Read gpio pin value from sysfs. ....	6
<b>2. Module debugging (gpio_input.c) .....</b>	7
Step 1: Compile module with debugging symbols.....	7
Step 2: Copy module source code .ko files in to host.....	7
Step 3: Identify Section addresses.....	7
Step 4: Assign a breakpoint in Kernel .....	7
Step 5: Load module symbols .....	7
Step 6: read gpio_input device file from host side than automatically breakpoint comes to target side. ....	8
<b>3. Linux Kernel Crash dump .....</b>	10
kexec: .....	10
kdump: .....	10
Enable the below kernel configuration options & compile: .....	10
Dump-capture kernel config options (Arch Independent): .....	11
Dump-capture kernel config options (Arch Dependent, i386 and x86_64).....	11
Dump-capture kernel config options (Arch Dependent, arm).....	12
crash kernel recommendations: .....	12
<b>Example:.....</b>	12
Core Dump Analyse:.....	14
1. using gdb tool.....	14
2. Using crash tool.....	15

3. apport-retrace tool: .....18

**Revision History**

1.0	Basic Version created By Kishore kumar Boddu

## 1. KDB & KGDB

### **KDB- Kernel Debugger:**

- Kdb is an instruction-level debugger used for debugging kernel code and device drivers.
- Before you can use it, you need to enable configuration options your kernel sources with kdb support and recompile the kernel.

### **KGDB- Kernel GNU Debugger:**

- Kgdb is a source-level debugger.
- It is easier to use than kdb because you don't have to spend time correlating assembly code with your sources.
- However it's more difficult to set up because an additional machine is needed to front-end the debugging.
- gdb runs on the host machine, while the kgdb-patched kernel runs on the target hardware.
- The host and the target are connected via a serial null-modem cable.

### **Setup Kernel Debug Environment for KDB & KGDB**

#### **Step 1: Enable KGDB/KDB options**

```
$ make ARCH=arm CROSS_COMPILE=arm-linux-gnueabihf- menuconfig
```

CONFIG\_DEBUG\_INFO: Compiles the kernel with debug symbols (-g option)

CONFIG\_DEBUG\_INFO=y

---> Kernel Hacking

    --> Compile-time checks and compiler options

    --> Compile the kernel with debug info

CONFIG\_KGDB: enables support for KGDB

CONFIG\_KGDB=y

---> Kernel Hacking

    ---> KGDB: Kernel Debugger

CONFIG\_KGDB\_SERIAL\_CONSOLE: Enables KGDB communication I/O driver over the serial port

CONFIG\_KGDB\_SERIAL\_CONSOLE=y

---> Kernel Hacking

    ---> KGDB: Kernel Debugger

    ---> KGDB: use kgdb over the serial console

CONFIG\_KGDB\_KDB=y

---> Kernel Hacking

    ---> KGDB: Kernel Debugger

    ---> KGDB\_KDB: include kdb frontend for kgdb

CONFIG\_KDB\_KEYBOARD=y

---> Kernel Hacking

    ---> KGDB: Kernel Debugger

    ---> KGDB\_KDB: include kdb frontend for kgdb

CONFIG\_FRAME\_POINTER: Helps to produce more reliable stack traces

**CONFIG\_MAGIC\_SYSRQ**: Enables magic sysrq key functionality to put the kernel in debug mode

---> Kernel Hacking

    ---> KGDB: Magic SysRq key

**CONFIG\_DEBUG\_RODATA\_TEST not set**

    ---> Kernel Hacking

        ---> Memory Debugging

            ---> Testcase for the marking rodata read-only (disable this option)

**CONFIG\_STRICT\_KERNEL\_RWX is not set**

**Step 2: Kernel configuration & build**

    \$ ./km-bbb-kernel-build.sh

**Step 3: Kernel Installation**

    \$ ./km-bbb-kernel-install.sh

**Step 4: Setup kernel bootargs in u-boot**

Boot any one of the mode MMC0, MMC1 & TFTP.

```
setenv bootargs console=tty0 console=${console} root=/dev/mmcblk1p1  
rootfstype=${mmcrootfstype} ${cmdline}; root=/dev/mmcblk1p1 console=ttyO0,115200n8  
rodata=off nokaslr kgdb=ttyO0,115200 kgdboc=ttyO0,115200n8 kgdbwait rootwait;
```

Kernel Address Space Randomization (KASLR)

```
setenv bootargs rootfstype=${mmcrootfstype} root=/dev/mmcblk0p1  
console=ttyS0,115200n8 rodata=off nokaslr kgdb=ttyS0,115200  
kgdboc=ttyS0,115200n8 kgdbwait rootwait;
```

**KDB Testing:**

**Choose KDB/KGDB Enabled Kernel at bootloader:**

Power on the board, choose "Boot with KGDB/KDB" option in "KM boot menu" at u-boot prompt.

By default kernel enter to KDB mode.

Kernel assign a break point after serial initialization and wait for user input.

```
kdb> help (list of kdb commands)  
kdb> go (continue kernel boot process)
```

Login as user until you see prompt.

**Entering a Kernel Debugger:**

*Case 1: User Entering a Kernel Debugger [KDB/KGDB]:*

Enter supervisor mode.

\$ sudo su

assign a break point to kernel.

\$ echo g > /proc/sysrq-trigger

Kernel enter to kdgb break point and shows kdb prompt.

kdb>

*Case 2: Whenever Kernel panic is occurs by default Enter KDB prompt.*

## kdb commands

[0]kdb> help	Command	Usage	Description
md <vaddr>		Display Memory Contents, also mdWcN, e.g1	
mdr <vaddr> <bytes>		Display Raw Memory	
mdp <paddr> <bytes>		Display Physical Memory	
mds <vaddr>		Display Memory Symbolically	
mm <vaddr> <contents>		Modify Memory Contents	
go [<vaddr>]		Continue Execution	
rd		Display Registers	
rm <reg> <contents>		Modify Registers	
ef <vaddr>		Display exception frame	
bt [<vaddr>]		Stack traceback	
btp <pid>		Display stack for process <pid>	
bta [D R S T C Z E U I M A]		Backtrace all processes matching state fg	
btc		Backtrace current process on each cpu	
btt <vaddr>		Backtrace process given its struct task s	
env		Show environment variables	
set		Set environment variables	
help		Display Help Message	
?		Display Help Message	
cpu <cpunum>		Switch to new cpu	
kgdb		Enter kgdb mode	
ps [<flags> A]		Display active task list	
pid <pidnum>		Switch to another task	
reboot		Reboot the machine immediately	
lsmod		List loaded kernel modules	
sr <key>		Magic SysRq key	
dmesg [lines]		Display syslog buffer	
defcmd name "usage" "help"		Define a set of commands, down to endefcd	
kill <-signal> <pid>		Send a signal to a process	
summary		Summarize the system	
per_cpu <sym> [<bytes>] [<cpu>]		Display per_cpu variables	
grepelp		Display help on   grep	
bp [<vaddr>]		Set/Display breakpoints	
bl [<vaddr>]		Display breakpoints	
bc <bpnum>		Clear Breakpoint	
be <bpnum>		Enable Breakpoint	
bd <bpnum>		Disable Breakpoint	
ss		Single Step	
dumpcommon		Common kdb debugging	
dumpall		First line debugging	
dumpcpu		Same as dumpall but only tasks on cpus	
ftdump [skip_#lines] [cpu]		Dump ftrace log	

### KGDB Debugging (source level debugging):

Whenever kernel crash, panic occurs by default enter in to KDB prompt. Switch from KDB to KGDB.

#### **Step 1: Connect target board remote using ssh.**

Run the below command in host machine and connect target board remote. Provide target board user name and IP address to the below command.

```
$ ssh <user_name>@<IP_Address>
```

Once remote login target board and run the below command to assign a break point to the kernel.

```
$ sudo su  
# echo g > /proc/sysrq-trigger
```

#### **Step 2: Switch KDB to KGDB**

Enter kgdb command in kdb prompt.

```
kdb> kgdb <ENTER>
```

You will now see "Entering please attach debugger or use \$D#44 or \$3#33"  
Close serial port terminal (minicom or teraterm) and open terminal in host machine.  
Assuming serial port device name is ttyUSBO

#### **Step 3: Remote Debugging with GDB**

**\$arm-linux-gnueabihf-gdb vmlinl** (this command shows gdb prompt it is called KGDB because image name is “vmlinux”)

```
(gdb) set serial baud 115200  
(gdb) target remote /dev/ttyUSBO  
(gdb) bt (shows back trace of kgdb_breakpoint())  
(gdb) b omap_gpio_get  
(gdb) c (continue)
```

#### **Step 4: Read gpio pin value from sysfs.**

```
$ cd /sys/class/gpio/gpio9  
$ cat value
```

Than you will see the gdb prompt with break point triggered

## 2. Module debugging (gpio\_input.c)

**Step 1: Compile module with debugging symbols**

Add KBUILD\_CFLAGS += -g flag in top most Makefile in Kernel source code.

**Step 2: Copy module source code .ko files in to host**

```
$ scp char.ko <username>@<IPAddress>:~/  
$ sudo insmod char.ko
```

**Step 3: Identify Section addresses**

```
$ sudo su // Should be in root user  
# cd /sys/modules/char/sections  
# ls -la  
# cat .text .data .bss
```

**Step 4: Assign a breakpoint in Kernel**

**Host Machine:**

```
$ sudo su  
# echo g > /proc/sysrq-trigger
```

**Step 5: Load module symbols**

**gdb prompt:**

```
(gdb) add-symbol-file /home/km/debug/kernel/gpio/gpio-input.ko 0xbff428000 -s .data 0xbff42a000 -s .bss 0xbff42a300
```

```
add symbol table from file "/home/km/debug/kernel/gpio/gpio-input.ko" at  
.text_addr = 0xbff428000  
.data_addr = 0xbff42a000  
.bss_addr = 0xbff42a300
```

(y or n) y

Reading symbols from /home/km/debug/kernel/gpio/gpio-input.ko...done.

```
(gdb) disass example_read
```

Dump of assembler code for function example\_read:

```
0xbff428000 <+0>: push {r4, r5, r6, r7, r8, lr}  
0xbff428004 <+4>: sub sp, sp, #40 ; 0x28  
0xbff428008 <+8>: ldr r7, [pc, #216] ; 0xbff4280e8 <example_read+232>  
0xbff42800c <+12>: ldr r0, [pc, #216] ; 0xbff4280ec <example_read+236>  
0xbff428010 <+16>: mov r4, r2  
0xbff428014 <+20>: ldr r3, [r7]  
0xbff428018 <+24>: ldr r0, [r0]  
0xbff42801c <+28>: mov r6, r1  
0xbff428020 <+32>: str r3, [sp, #36] ; 0x24  
0xbff428024 <+36>: bl 0xc0460d88 <gpio_to_desc>  
0xbff428028 <+40>: bl 0xc045fa90 <gpiod_get_raw_value>  
0xbff42802c <+44>: ldr r2, [pc, #188] ; 0xbff4280f0 <example_read+240>  
0xbff428030 <+48>: mov r1, #32  
0xbff428034 <+52>: mov r3, r0  
0xbff428038 <+56>: add r0, sp, #4
```

```

0xbf42803c <+60>: bl      0xc081e340 <snprintf>
--Type <RET> for more, q to quit, c to continue without paging--q
Quit

(gdb) b example_read
Breakpoint 1 at 0xbf428000: file /home/km/debug/kernel/gpio/gpio-input.c, line 87.
(gdb) b omap_gpio_get
Breakpoint 2 at 0xc046702c: file drivers/gpio/gpio-omap.c, line 1009.
(gdb) c
Continuing.

```

Step 6: read gpio\_input device file from host side than automatically breakpoint comes to target side.

**Host Machine:**

```

$ sudo su
# cat /dev/gpio_input

```

Thread 96 hit Breakpoint 1, example\_read (filp=0xdb7727c0, buffer=0xb6982000 <error: Cannot access memory at address 0xb6982000>, length=131072, offset=0xdb75ff78)

```

    at /home/km/debug/kernel/gpio/gpio-input.c:87
87      {

```

```

(gdb) i b
Num  Type      Disp Enb Address  What
1    breakpoint  keep y  0xbf428000 in example_read at /home/km/debug/kernel/gpio/gpio-
input.c:87
      breakpoint already hit 1 time
2    breakpoint  keep y  0xc046702c in omap_gpio_get at drivers/gpio/gpio-omap.c:1009

(gdb) bt
#0 example_read (filp=0xdb7727c0, buffer=0xb6982000 <error: Cannot access memory at address
0xb6982000>, length=131072, offset=0xdb75ff78)
    at /home/km/debug/kernel/gpio/gpio-input.c:87
#1 0xc01f4348 in vfs_read (pos=<optimized out>, count=<optimized out>, buf=<optimized out>,
file=<optimized out>) at fs/read_write.c:452
#2 vfs_read (file=0xdb7727c0, buf=0xb6982000 <error: Cannot access memory at address
0xb6982000>, count=131072, pos=0xdb75ff78) at fs/read_write.c:437
#3 0xc01f4748 in ksys_read (fd=<optimized out>, buf=0xb6982000 <error: Cannot access memory at
address 0xb6982000>, count=131072) at fs/read_write.c:605
#4 0xc0009000 in __idmap_text_end ()
Backtrace stopped: previous frame identical to this frame (corrupt stack?)


```

Thread 96 hit Breakpoint 1, example\_read (filp=0xdb7727c0, buffer=0xb6982000 "1",
length=131072, offset=0xdb75ff78) at /home/km/debug/kernel/gpio/gpio-input.c:87

```

87      {
(gdb) c
Continuing.

```

```
Thread 96 hit Breakpoint 2, omap_gpio_get (chip=0xde18ecd4, offset=11) at drivers/gpio/gpio-omap.c:1009
```

```
1009          bank = gpiochip_get_data(chip);
(gdb) c
Continuing.
```

```
Thread 96 hit Breakpoint 1, example_read (filp=0xdb7727c0, buffer=0xb6982000 "1",
length=131072, offset=0xdb75ff78) at /home/km/debug/kernel/gpio/gpio-input.c:87
```

```
87      {
(gdb) c
Continuing.
```

```
Thread 96 hit Breakpoint 2, omap_gpio_get (chip=0xde18ecd4, offset=11) at drivers/gpio/gpio-omap.c:1009
```

```
1009          bank = gpiochip_get_data(chip);
(gdb) c
Continuing.
```

```
Thread 96 hit Breakpoint 1, example_read (filp=0xdb7727c0, buffer=0xb6982000 "1",
length=131072, offset=0xdb75ff78) at /home/km/debug/kernel/gpio/gpio-input.c:87
```

```
87      {
(gdb) n
80      device_destroy(example_class, example_dev);
(gdb) n
81      class_destroy(example_class);
(gdb) n
90      sprintf(chaine, 32, "%d", gpio_get_value(gpio_in));
(gdb) n
```

```
Thread 96 hit Breakpoint 2, omap_gpio_get (chip=0xde18ecd4, offset=11) at drivers/gpio/gpio-omap.c:1009
```

```
1009          bank = gpiochip_get_data(chip);
(gdb) n
1011          if (omap_gpio_is_input(bank, offset))
(gdb) n
1012              return omap_get_gpio_datain(bank, offset);
(gdb) n
1014          return omap_get_gpio_dataout(bank, offset);
(gdb) n
gpiod_get_raw_value_commit (desc=0xde18eab0) at drivers/gpio/gpiolib.c:2832
2832          value = value < 0 ? value : !value;
(gdb) p value
$10 = <optimized out>
(gdb) n
2833          trace_gpio_value(desc_to_gpio(desc), 1, value);
(gdb) n
warning: Error removing breakpoint 0
[New Thread 752]
```

### 3. Linux Kernel Crash dump

Install crash dump tools:

```
$ sudo apt install kexec-tools
```

```
$ sudo apt-get install kdump-tools
```

Using kdump and kexec analyze the kernel coredump files.

#### **kexec:**

kexec is directly boot into a new kernel without going through BIOS.

Kexec is a fastboot mechanism that allows booting a Linux kernel from the context of an already running kernel without going through BIOS.

BIOS can be very time consuming, especially on big servers with numerous peripherals.

This can save a lot of time for developers who end up booting a machine numerous times **kexec is a system call that enables you to load and boot into another kernel** from the currently running kernel. kexec performs the function of the boot loader from within the kernel.

The primary difference between a standard system boot and a kexec boot is that the hardware initialization normally performed by the BIOS or firmware (depending on architecture) is not performed during a kexec boot. This has the effect of reducing the time required for a reboot

#### **kdump:**

Kdump is a new kernel crash dumping mechanism and is very reliable.

The crash dump is captured from the context of a freshly booted kernel and not from the context of the crashed kernel.

Kdump uses Kexec to boot into a second kernel whenever the system crashes.

This second kernel, often called a crash or a capture kernel, boots with very little memory and captures the dump image.

The first kernel reserves a section of memory that the second kernel uses to boot.

Kexec enables booting the capture kernel without going through BIOS hence the contents of the first kernel's memory are preserved, which is essentially the kernel crash dump.

Kdump is a feature of the Linux kernel that creates crash dumps in the event of a kernel crash.

When triggered, kdump exports a memory image (also known as vmcore) that can be analyzed for the purposes of debugging and determining the cause of a crash.

#### **Enable the below kernel configuration options & compile:**

System kernel config options:

1) Enable "kexec system call" in "Processor type and features."

This parameter tells the system to use Kexec to skip BIOS and boot (new) kernels. It is critical for the functionality of Kdump.

```
CONFIG_KEXEC=y
```

```
Boot options -->
```

```
[*] Kexec system call (EXPERIMENTAL)
```

2) Enable "sysfs file system support" in "Filesystem" -> "Pseudo filesystems." This is usually enabled by default

```
CONFIG_SYSFS=y
```

3) Enable "Compile the kernel with debug info" in "Kernel hacking."

This causes the kernel to be built with debug symbols. The dump analysis tools require a vmlinu with debug symbols in order to read and analyze a dump file

CONFIG\_DEBUG\_INFO=y

### Dump-capture kernel config options (Arch Independent):

1) Enable kernel crash dumps: Crash dumps need to be enabled. Without this option, Kdump will be useless.

CONFIG\_CRASH\_DUMP=y

Boot options --->

Build kdump crash kernel (EXPERIMENTAL)

2) Enable "/proc/vmcore support" under "Filesystems" -> "Pseudo filesystems".

CONFIG\_PROC\_VMCORE=y

(CONFIG\_PROC\_VMCORE is set by default when CONFIG\_CRASH\_DUMP is selected.)

-> File systems

-> Pseudo filesystems

-> /proc file system support (PROC\_FS [=y])

### Dump-capture kernel config options (Arch Dependent, i386 and x86\_64)

1) If one wants to build and use a relocatable kernel,

Enable "Build a relocatable kernel" support under "Processor type and features"

CONFIG\_RELOCATABLE=y

2) Optional: Disable Symmetric Multi-Processing (SMP) support: Kdump can only work with a single processor.

If you have only a single processor or you run your machine with SMP support disabled, you can safely set this parameter to (n).

CONFIG\_SMP=n

KDUMP\_COMMANDLINE\_APPEND="maxcpus=1"

(If CONFIG\_SMP=y, then specify maxcpus=1 on the kernel command line when loading the dump-capture kernel, see section "Load the Dump-capture Kernel".)

3) Use a suitable value for "Physical address where the kernel is loaded" (under "Processor type and features").

This only appears when "kernel crash dumps" is enabled. A suitable value depends upon whether kernel is relocatable or not.

If you are using a relocatable kernel use CONFIG\_PHYSICAL\_START=0x100000

This will compile the kernel for physical address 1MB, but given the fact kernel is relocatable, it can be run from any physical address hence kexec boot loader will load it in memory region reserved for dump-capture kernel.

Otherwise it should be the start of memory region reserved for second kernel using boot parameter "crashkernel=Y@X".

Here X is start of memory region reserved for dump-capture kernel.

Generally X is 16MB (0x1000000). So you can set CONFIG\_PHYSICAL\_START=0x1000000

5) Make and install the kernel and its modules.

## Dump-capture kernel config options (Arch Dependent, arm)

1) To use a relocatable kernel,

Enable "AUTO\_ZRELADDR" support under "Boot" options:

AUTO\_ZRELADDR=y

### crash kernel recommendations:

For memory between 2G through 4G, reserve 320M

For memory between 4G through 32G, reserve 512M

For memory between 32G through 64G, reserve 1024M

For memory between 64G through 128G, reserve 2048M

For memory above 128G, reserve 4096M

add "crashkernel=256M" in /boot/grub/grub.cfg file.

### Example:

Command line: BOOT\_IMAGE=/boot/vmlinuz-4.4.88 root=UUID=aab9b05c-202b-42c6-8be6-c90e9d6af8e8 ro quiet splash crashkernel=256M vt.handoff=7

```
$ kdump-config status
```

```
$ kdump-config show
```

```
$ kdump-config load
```

```
$ cat /proc/sys/kernel/sysrq
```

```
$ sudo sysctl -w kernel.sysrq=1
```

```
$ sudo kexec -l /boot/vmlinuz-4.19.132 --append="BOOT_IMAGE=/boot/vmlinuz-4.19.132
root=UUID=f51c4fff-68ae-4bb9-86dd-5d8716be2821 ro quiet splash" --initrd=/boot/initrd.img-4.19.132
```

```
$ sudo kexec -l /boot/vmlinuz-4.19.94-Kernel-Masters-g001db1705-dirty --
append="console=ttyO0,115200n8 root=/dev/mmcblk1p1 rootfstype=ext4 rootwait"
```

```
$ sudo kexec -e
```

```
kexec --type zImage -p /boot/vmlinuz-4.19.94-Kernel-Masters-gb3f1becdc-dirty --
initrd=/var/lib/kdump/initrd.img-4.19.94-Kernel-Masters-g001db1705-dirty --
dtb=/boot/dtbs/4.19.94-Kernel-Masters-gb3f1becdc-dirty/am335x-boneblack.dtb --
append="console=ttyO0,115200n8 root=/dev/mmcblk1p1 rootfstype=ext4 rootwait
crashkernel=64B maxcpus=1 reset_devices"
```

```
kexec --type zImage -p --command-line="console=ttyO0,115200n8 root=/dev/mmcblk1p1
rootfstype=ext4 rootwait crashkernel=64B maxcpus=1 reset_devices" /boot/vmlinuz-4.19.94-Kernel-
Masters-gb3f1becdc-dirty --initrd=/var/lib/kdump/initrd.img-4.19.94-Kernel-Masters-g001db1705-
dirty --dtb=/boot/dtbs/4.19.94-Kernel-Masters-gb3f1becdc-dirty/am335x-boneblack.dtb
```

```
/sbin/kexec -u /boot/vmlinuz-4.19.94-Kernel-Masters-gb3f1becdc-dirty --dtb=/boot/dtbs/4.19.94-
Kernel-Masters-gb3f1becdc-dirty/am335x-boneblack.dtb --command-line="console=ttyO0,115200n8
root=/dev/mmcblk1p1 rootfstype=ext4 rootwait maxcpus=1 reset_devices nr_cpus=1
systemd.unit=kdump-tools.service irqpoll nousb ata_piix.prefer_ms_hyperv=0" --
initrd=/var/lib/kdump/initrd.img /var/lib/kdump/vmlinuz
```

```
kexec --type=zImage -p /boot/vmlinuz-4.19.94-Kernel-Masters-gb3f1becdc-dirty --
dtb=/boot/dtbs/4.19.94-Kernel-Masters-gb3f1becdc-dirty/am335x-boneblack.dtb --
append="console=ttyO0,115200n8 root=/dev/mmcblk1p1 rootfstype=ext4 rootwait maxcpus=1
reset_devices"
```

Usage: kexec [OPTION]... [kernel]

Directly reboot into a new kernel

- h, --help Print this help.
- v, --version Print the version of kexec.
- f, --force Force an immediate kexec,  
don't call shutdown.
- i, --no-checks Fast reboot, no memory integrity checks.
- x, --no-ifdown Don't bring down network interfaces.
- y, --no-sync Don't sync filesystems before kexec.
- l, --load Load the new kernel into the  
current kernel.
- p, --load-panic Load the new kernel for use on panic.
- u, --unload Unload the current kexec target kernel.  
If capture kernel is being unloaded  
specify -p with -u.
- e, --exec Execute a currently loaded kernel.
- t, --type=TYPE Specify the new kernel is of this type.
- mem-min=<addr> Specify the lowest memory address to  
load code into.
- mem-max=<addr> Specify the highest memory address to  
load code into.
- reuseinitrd Reuse initrd from first boot.
- print-ckr-size Print crash kernel region size.
- load-preserve-context Load the new kernel and preserve  
context of current kernel during kexec.
- load-jump-back-helper Load a helper image to jump back  
to original kernel.
- entry=<addr> Specify jump back address.  
(0 means it's not jump back or  
preserve context)  
to original kernel.
- s, --kexec-file-syscall Use file based syscall for kexec operation
- c, --kexec-syscall Use the kexec\_load syscall for compatibility  
with systems that don't support -s (default)
- a, --kexec-syscall-auto Use file based syscall for kexec and fall  
back to the compatibility syscall when file based  
syscall is not supported or the kernel did not  
understand the image
- d, --debug Enable debugging to help spot a failure.
- S, --status Return 0 if the type (by default crash) is loaded.

Supported kernel file types and options:

ulimage

- command-line=STRING Set the kernel command line to STRING.
- append=STRING Set the kernel command line to STRING.

```

--initrd=FILE      Use FILE as the kernel's initial ramdisk.
--ramdisk=FILE     Use FILE as the kernel's initial ramdisk.
--dtb=FILE        Use FILE as the fdt blob.
--atags          Use ATAGs instead of device-tree.
--page-offset=PAGE_OFFSET
                  Set PAGE_OFFSET of crash dump vmcore

zImage
--command-line=STRING Set the kernel command line to STRING.
--append=STRING     Set the kernel command line to STRING.
--initrd=FILE      Use FILE as the kernel's initial ramdisk.
--ramdisk=FILE     Use FILE as the kernel's initial ramdisk.
--dtb=FILE        Use FILE as the fdt blob.
--atags          Use ATAGs instead of device-tree.
--page-offset=PAGE_OFFSET
                  Set PAGE_OFFSET of crash dump vmcore

Architecture options:
--image-size=<size>
                  Specify the assumed total image size of
                  the kernel that is about to be loaded,
                  including the .bss section, as reported
                  by 'arm-linux-size vmlinux'. If not
                  specified, this value is implicitly set
                  to the compressed images size * 4.
--dt-no-old-root
                  do not reuse old kernel root= param.
                  while creating flatten device tree.

```

## Core Dump Analyse:

### 1. using gdb tool

```

$ sudo gdb KM_GIT/linux-4.4.88/vmlinux /var/crash/201712022302/vmcore.201712022302
[sudo] password for kernel:
GNU gdb (Ubuntu 7.7.1-0ubuntu5~14.04.3) 7.7.1
Copyright (C) 2014 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law. Type "show copying"
and "show warranty" for details.
This GDB was configured as "x86_64-linux-gnu".
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
<http://www.gnu.org/software/gdb/documentation/>.
For help, type "help".
Type "apropos word" to search for commands related to "word"...
Reading symbols from KM_GIT/linux-4.4.88/vmlinux...done.
[New LWP 2451]
[New LWP 1120]
[New process 1]

```

```

[New LWP 2142]
#0 sysrq_handle_crash (key=99) at drivers/tty/sysrq.c:138
138          *killer = 1;
warning: File "/home/kernel/KM_GIT/linux-4.4.88/scripts/gdb/vmlinux-gdb.py" auto-loading has
been declined by your `auto-load safe-path' set to "$debugdir:$datadir/auto-load".
To enable execution of this file add
    add-auto-load-safe-path /home/kernel/KM_GIT/linux-4.4.88/scripts/gdb/vmlinux-gdb.py
line to your configuration file "/home/kernel/.gdbinit".
To completely disable this security protection add
    set auto-load safe-path /
line to your configuration file "/home/kernel/.gdbinit".
For more information about this security protection see the
"Auto-loading safe path" section in the GDB manual. E.g., run from the shell:
    info "(gdb)Auto-loading safe path"
(gdb) bt
#0 sysrq_handle_crash (key=99) at drivers/tty/sysrq.c:138
#1 0xffffffff814c28c9 in __handle_sysrq (key=99, check_mask=false) at drivers/tty/sysrq.c:545
#2 0xffffffff814c2d18 in write_sysrq_trigger (file=<optimized out>, buf=<optimized out>, count=1,
ppos=<optimized out>)
    at drivers/tty/sysrq.c:1091
#3 0xffffffff8126828d in proc_reg_write (file=<optimized out>, buf=<optimized out>,
count=<optimized out>,
    ppos=<optimized out>) at fs/proc/inode.c:216
#4 0xffffffff811fd558 in __vfs_write (file=0xffff8802411b8700, p=<optimized out>,
count=<optimized out>,
    pos=0x0 <irq_stack_union>) at fs/read_write.c:489
#5 0xffffffff811fdc12 in vfs_write (file=0x63 <irq_stack_union+99>, buf=0xffff88024f40dd98 "",,
count=<optimized out>,
    pos=0xffff88024307bf20) at fs/read_write.c:538
#6 0xffffffff811fe936 in SYSC_write (count=<optimized out>, buf=<optimized out>, fd=<optimized
out>)
    at fs/read_write.c:585
#7 SyS_write (fd=<optimized out>, buf=15049736, count=2) at fs/read_write.c:577
#8 0xffffffff817e6e36 in entry_SYSCALL_64_fastpath () at arch/x86/entry/entry_64.S:185
#9 0x00000000000e8f608 in ?? ()
#10 0x0000000000000004 in irq_stack_union ()

```

## 2. Using crash tool

```

$ cd /var/crash/
$ ls
201902261006 kexec_cmd
$ cd 201902261006/

$ sudo apt source linux-image-`uname -r` 

$ crash <debug kernel> <crash dump>

$ sudo crash KM_GIT/linux-4.4.88/vmlinux /var/crash/201712022302/vmcore.201712022302

crash 7.0.3

```

Copyright (C) 2002-2013 Red Hat, Inc.  
Copyright (C) 2004, 2005, 2006, 2010 IBM Corporation  
Copyright (C) 1999-2006 Hewlett-Packard Co  
Copyright (C) 2005, 2006, 2011, 2012 Fujitsu Limited  
Copyright (C) 2006, 2007 VA Linux Systems Japan K.K.  
Copyright (C) 2005, 2011 NEC Corporation  
Copyright (C) 1999, 2002, 2007 Silicon Graphics, Inc.  
Copyright (C) 1999, 2000, 2001, 2002 Mission Critical Linux, Inc.  
This program is free software, covered by the GNU General Public License,  
and you are welcome to change it and/or distribute copies of it under  
certain conditions. Enter "help copying" to see the conditions.  
This program has absolutely no warranty. Enter "help warranty" for details.

GNU gdb (GDB) 7.6  
Copyright (C) 2013 Free Software Foundation, Inc.  
License GPLv3+: GNU GPL version 3 or later <<http://gnu.org/licenses/gpl.html>>  
This is free software: you are free to change and redistribute it.  
There is NO WARRANTY, to the extent permitted by law. Type "show copying"  
and "show warranty" for details.  
This GDB was configured as "x86\_64-unknown-linux-gnu"...

KERNEL: KM\_GIT/linux-4.4.88/vmlinux  
DUMPFILE: /var/crash/201712022302/vmcore.201712022302  
CPUS: 4  
DATE: Thu Jan 1 05:30:00 1970  
UPTIME: 00:05:51  
LOAD AVERAGE: 0.36, 0.52, 0.29  
TASKS: 367  
NODENAME: KernelMasters  
RELEASE: 4.4.88  
VERSION: #1 SMP Fri Dec 1 22:09:18 IST 2017  
MACHINE: x86\_64 (2711 Mhz)  
MEMORY: 7.9 GB  
PANIC: "Oops: 0002 [#1] SMP " (check log for details)  
PID: 2451  
COMMAND: "bash"  
TASK: ffff880243510000 [THREAD\_INFO: ffff880243078000]  
CPU: 0  
STATE: TASK\_RUNNING (PANIC)

crash>  
crash> help

*	files	mach	repeat	timer
alias	foreach	mod	runq	tree
ascii	fuser	mount	search	union
bt	gdb	net	set	vm
btop	help	p	sig	vtop
dev	ipcs	ps	struct	waitq
dis	irq	pte	swap	whatis

```
eval      kmem      ptob      sym      wr
exit      list      ptov      sys      q
extend    log       rd       task
```

crash version: 7.0.3 gdb version: 7.6

For help on any command above, enter "help <command>".

For help on input options, enter "help input".

For help on output options, enter "help output".

crash> ps

PID	PPID	CPU	TASK	ST	%MEM	VSZ	RSS	COMM
0	0	0	ffffffffff81c13500	RU	0.0	0	0	[swapper/0]
0	0	1	ffff880245699e00	RU	0.0	0	0	[swapper/1]
0	0	2	ffff88024569ad00	RU	0.0	0	[swapper/2]	
0	0	3	ffff88024569bc00	RU	0.0	0	0	[swapper/3]
1	0	3	ffff880245b58000	IN	0.0	34060	4576	init

crash> bt

PID: 2451 TASK: ffff880243510000 CPU: 0 COMMAND: "bash"  
#0 [ffff88024307bac0] machine\_kexec at ffffffff8105919c  
#1 [ffff88024307bb18] crash\_kexec at ffffffff81106153  
#2 [ffff88024307bbe0] oops\_end at ffffffff8101aae9  
#3 [ffff88024307bc08] no\_context at ffffffff81066f3d  
#4 [ffff88024307bc60] \_\_bad\_area\_nosemaphore at ffffffff810672b9  
#5 [ffff88024307bca8] bad\_area\_nosemaphore at ffffffff810673d3  
#6 [ffff88024307bcb8] \_\_do\_page\_fault at ffffffff810679c0  
#7 [ffff88024307bd10] do\_page\_fault at ffffffff81067d82  
#8 [ffff88024307bd30] page\_fault at ffffffff817e8fb8  
[exception RIP: sysrq\_handle\_crash+22]  
RIP: ffffffff814c2116 RSP: ffff88024307bde0 RFLAGS: 00010296  
RAX: 0000000000000000 RBX: ffffffff81cbb800 RCX: 0000000000000000  
RDX: 0000000000000001 RSI: ffff88024f40dd98 RDI: 0000000000000063  
RBP: ffff88024307bde0 R8: ffffffff81f19ed4 R9: 0000000000000030  
R10: ffffffff81f0937c R11: 0000000000000358 R12: 0000000000000063  
R13: 0000000000000000 R14: 0000000000000004 R15: 0000000000000000  
ORIG\_RAX: ffffffffffffd CS: 0010 SS: 0018  
#9 [ffff88024307bde8] \_\_handle\_sysrq at ffffffff814c28c9  
#10 [ffff88024307be18] write\_sysrq\_trigger at ffffffff814c2d18  
#11 [ffff88024307be30] proc\_reg\_write at ffffffff8126828d  
#12 [ffff88024307be50] \_\_vfs\_write at ffffffff811fd558  
#13 [ffff88024307bed0] vfs\_write at ffffffff811fdc12  
#14 [ffff88024307bf10] sys\_write at ffffffff811fe936  
#15 [ffff88024307bf50] entry\_SYSCALL\_64\_fastpath at ffffffff817e6e36  
RIP: 00007f6f67b15390 RSP: 00007ffca2f8acd8 RFLAGS: 00000246  
RAX: ffffffffffffd RBX: 0000000004b7ddd RCX: 00007f6f67b15390  
RDX: 0000000000000002 RSI: 0000000000e5a408 RDI: 0000000000000001  
RBP: 00000000004b7ddd R8: 000000000000000a R9: 00007f6f68423740  
R10: 00007f6f67de76a0 R11: 0000000000000246 R12: 0000000000e74e88  
R13: 0000000000000004 R14: 0000000000000004 R15: 0000000000e8f608  
ORIG\_RAX: 0000000000000001 CS: 0033 SS: 002b

crash>

```

crash>
crash> kmem -i
      PAGES   TOTAL   PERCENTAGE
TOTAL MEM 1944950    7.4 GB    ---
  FREE 1641126    6.3 GB  84% of TOTAL MEM
 USED 303824     1.2 GB 15% of TOTAL MEM
SHARED 43821    171.2 MB  2% of TOTAL MEM
BUFFERS 15850    61.9 MB  0% of TOTAL MEM
CACHED 136661    533.8 MB  7% of TOTAL MEM
 SLAB 13080    51.1 MB  0% of TOTAL MEM

TOTAL SWAP 3999999    15.3 GB    ---
SWAP USED     0       0  0% of TOTAL SWAP
SWAP FREE 3999999    15.3 GB 100% of TOTAL SWAP

crash> irq
IRQ IRQ_DESC/_DATA   IRQACTION   NAME
 0 ffff88024682e800 ffffffff81c19a40 "timer"
 1 ffff88024682ea00 ffff88022c002f80 "i8042"
 2 ffff88024682ec00  (unused)
 3 ffff88024682ee00  (unused)
 4 ffff88024682f000  (unused)
 5 ffff88024682f200  (unused)
 6 ffff88024682f400  (unused)
 7 ffff88024682f600  (unused)
 8 ffff88024682f800 ffff88022c05c180 "rtc0"
 9 ffff88024682fa00 ffff88024516c100 "acpi"
10 ffff88024682fc00  (unused)
11 ffff88024682fe00  (unused)
12 ffff880246888000 ffff88022c002f00 "i8042"
13 ffff880246888200  (unused)
14 ffff880246888400 ffff8800a2cc7b00 "INT344B:00"
15 ffff880246888600  (unused)
16 ffff880245363800 ffff88009906a380 "idma64.0"
          ffff8800a6d07b00 "i2c_designware.0"
17 ffff88022cdffe00 ffff88009906a980 "idma64.1"

```

### **3. apport-retrace tool:**

```
$ apport-retrace --stdout --rebuild-package-info /var/crash/linux-image*.crash
```